

Simulating survey studies

Carlo Knotz

Table of contents

1	Introduction	1
1.1	Setup	2
1.2	Simulating our fictional population	3
1.3	Taking a first sample	5
1.4	Taking a second sample	7
1.5	Comparing the samples	9
1.6	Scaling it up to 1000 samples	10
1.7	Then we do 10'000 surveys	12
1.8	The results in context	14
1.9	Comparing our results distribution (“sampling distribution”) to the Normal distribution	15
1.10	Now 10'000 new samples, but with a smaller sample size (200 instead of 1'000)	17

1 Introduction

This document shows you how you can simulate many, many different iterations of a fictional survey study – and how the different results are, collectively, normally distributed.

Important: Some of these simulations take a few minutes to run. Be patient.

1.1 Setup

These are the packages you'll need (use `install.packages()` to install anything that is missing):

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr     1.1.4     v readr     2.1.5
v forcats   1.0.0     v stringr   1.5.1
v ggplot2   4.0.0     v tibble    3.2.1
v lubridate 1.9.4     v tidyr    1.3.1
v purrr    1.2.0

-- Conflicts -----
x dplyr::filter() masks stats::filter()
x dplyr::lag()    masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
theme_set(theme_classic())
library(bst290)
library(ggpubr)
```

Next, we also record the start time with `Sys.time()` so that we can later check how long the entire simulation took:

```
start <- Sys.time()
```

1.2 Simulating our fictional population

Next, we simulate our fictional population: 5.5 mio. individuals and their levels of happiness on a 0-10 scale:


```
set.seed(42)

pop <- data.frame(idno = seq(1,5.5*10^6,1),
                    happy = sample(seq(0,10,1),
                                  size = 5.5*10^6,
                                  replace = T,
                                  prob = c(0.01,.01,.02,.03,.04,.09,.10,.13,.20,.29,0.05)))

popmean <- mean(pop$happy)

pop |>
  group_by(happy) |>
  summarize(obs = n()) |>
  mutate(share = obs/sum(obs)) |>
  ggplot(aes(x = happy, y = share)) +
  geom_bar(stat = "identity") +
  geom_text(aes(label = round(100*share, digits = 1)), vjust = -1) +
  geom_vline(xintercept = popmean, color = "orange") +
  scale_x_continuous(breaks = seq(0,10,1)) +
  scale_y_continuous(labels = scales::percent,
                     limits = c(0,.35)) +
  labs(y = "Percent", x = "How happy are you?",
       title = "Entire population of 5.5 mil. (simulated)",
       caption = paste0("Mean in population ('true value') = ",round(popmean, digits = 2))) +
  theme_minimal()

popvis
```



Note that we set a “seed” number so that we can always exactly reproduce the same results later.

1.3 Taking a first sample


We pretend that we are doing a social science survey by taking a random sample from the population with `slice_sample()` and calculate the sample mean:

```
set.seed(42)
sam1 <- pop |>
  slice_sample(n = 1000)


sam1mean <- mean(sam1$happy)

sam1 |>
  group_by(happy) |>
  summarize(obs = n()) |>
  mutate(share = obs/sum(obs)) |>
  ggplot(aes(x = happy, y = share)) +
  geom_bar(stat = "identity") +
  geom_text(aes(label = round(100*share, digits = 1)), vjust = -1) +
  geom_vline(xintercept = sam1mean, color = "orange") +
  scale_x_continuous(breaks = seq(0,10,1)) +
  scale_y_continuous(labels = scales::percent,
                     limits = c(0,.35)) +
  labs(y = "Percent", x = "How happy are you?",
       title = "Sample 1 (N = 1000; randomly selected)",
       caption = paste0("Mean in sample ('estimate') = ",round(sam1mean, digits = 2))) -> sam1vis

ggpubr::ggarrange(popvis,sam1vis, nrow = 2)
```



Mean in population ('true value') = 7.19



Mean in sample ('estimate') = 7.28

1.4 Taking a second sample


Then we take a second sample:

```
set.seed(17)
sam2 <- pop |>
  slice_sample(n = 1000)


sam2mean <- mean(sam2$happy)

sam2 |>
  group_by(happy) |>
  summarize(obs = n()) |>
  mutate(share = obs/sum(obs)) |>
  ggplot(aes(x = happy, y = share)) +
  geom_bar(stat = "identity") +
  geom_text(aes(label = round(100*share, digits = 1)), vjust = -1) +
  geom_vline(xintercept = sam2mean, color = "orange") +
  scale_x_continuous(breaks = seq(0,10,1)) +
  scale_y_continuous(labels = scales::percent,
                      limits = c(0,.35)) +
  labs(y = "Percent", x = "How happy are you?",
       title = "Sample 2 (N = 1000; randomly selected)",
       caption = paste0("Mean in sample ('estimate') = ",round(sam2mean, digits = 2))) -> sam2vis

ggpubr::ggarrange(popvis,sam2vis, nrow = 2)
```




Mean in population ('true value') = 7.19



Mean in sample ('estimate') = 7.27

1.5 Comparing the samples

```
ggpubr::ggarrange(sam1vis,sam2vis, nrow = 2)
```



1.6 Scaling it up to 1000 samples

We collect 998 new samples so that we get to 1'000 in total:

```
sampdist <- data.frame(sample = seq(1,1000,1),
                        result = c(c(sam1mean,sam2mean),rep(NA,998)))

for(k in 3:1000) {

  loopsam <- pop |>
    slice_sample(n = 1000)

  sampdist[k,2] <- mean(loopsam$happy)
}


sampdistmean <- mean(sampdist$result)

sampdist |>
  ggplot(aes(x = result)) +
  geom_histogram(color = "white") +
  geom_vline(xintercept = sampdistmean, color = "cornflowerblue", linetype = "dashed") +
  geom_vline(xintercept = popmean, color = "orange") +
  labs(x = "Sample means", y = "Samples",
       title = "Distribution of results ('sampling distribution')",
       caption = paste0("Population mean ('true value'; orange) = ",round(popmean, digits = 2),
                      "\n Mean over all results ('mean of means'; blue, dashed) = ",round(
```



```
`stat_bin()` using `bins = 30`. Pick better value `binwidth`.
```

Distribution of results ('sampling distribution')



1.7 Then we do 10'000 surveys

We're scaling it up once more to 10'000 surveys overall:

```
set.seed(42)
sampdist <- data.frame(sample = seq(1,10000,1),
                        result = c(c(sam1mean,sam2mean),rep(NA,9998)))


for(k in 3:10000) {

  loopsam <- pop |>
    slice_sample(n = 1000)

  sampdist[k,2] <- mean(loopsam$happy)
}

sampdistmean <- mean(sampdist$result)

sampdist |>
  ggplot(aes(x = result)) +
  geom_histogram(color = "white") +
  geom_vline(xintercept = sampdistmean, color = "cornflowerblue", linetype = "dashed") +
  geom_vline(xintercept = popmean, color = "orange") +
  labs(x = "Sample means", y = "Samples",
       title = "Distribution of results ('sampling distribution')",
       caption = paste0("Population mean ('true value'; orange) = ",round(popmean, digits = 2),
                       "\n Mean over all results ('mean of means'; blue, dashed) = ",round(sampdistmean, digits = 2)))
  `stat_bin()` using `bins = 30`. Pick better value `binwidth`.
```




1.8 The results in context

To see how much “wiggling” there really is due to random sampling variation:

```
sampdist |>
  ggplot(aes(x = result)) +
  geom_histogram(color = "white", binwidth = 0.05) +
  scale_x_continuous(limits = c(0,10), breaks = seq(0,10,1)) +
  labs(x = "Sample means", y = "Samples",
       title = "Distribution of results ('sampling distribution')")
```


Warning: Removed 2 rows containing missing values or values outside the scale range
(`geom_bar()`).



1.9 Comparing our results distribution (“sampling distribution”) to the Normal distribution


```
sampdist |>
  ggplot(aes(x = result)) +
  geom_histogram(color = "white") +
  geom_vline(xintercept = sampdistmean, color = "cornflowerblue", linetype = "dashed") +
  geom_vline(xintercept = popmean, color = "orange") +
  labs(x = "Sample means", y = "Samples",
       title = "Distribution of results ('sampling distribution')") -> samdist10k
samdist10k
```

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.




```
p1 <- ggplot(data = data.frame(x = c(-3, 3)), aes(x)) +
  stat_function(fun = dnorm, n = 101, args = list(mean = 0, sd = 1)) + ylab("") +
  scale_y_continuous(breaks = NULL) +
  scale_x_continuous(breaks = NULL) +
  labs(x = "", title = "Normal distribution")
p1
```

Normal distribution



```
ggpubr::ggarrange(samdist10k,p1, ncol = 2)
```

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.



1.10 Now 10'000 new samples, but with a smaller sample size (200 instead of 1'000)

We reduce the sample size to 200 and take 10'000 new samples:

```
sampdist_200 <- data.frame(sample = seq(1,10000,1),
                             result = rep(NA,10000))

set.seed(42)
for(k in 1:10000) {

  loopsam <- pop |>
    slice_sample(n = 200)


  sampdist_200[k,2] <- mean(loopsam$happy)
}

sampdist_200mean <- mean(sampdist_200$result)

sampdist_SD <- sd(sampdist$result)
sampdist_200_SD <- sd(sampdist_200$result)

sampdist |>
  left_join(sampdist_200, by = "sample",
            suffix = c("1000","200")) |>
  pivot_longer(cols = -1,
                names_to = "size",
                values_to = "meanval") |>
  mutate(size = gsub("result","",size)) |>
  ggplot(aes(x = meanval,fill = size)) +
  geom_histogram(alpha=0.3, position="identity",
                 color = "grey") +
  scale_fill_manual(values = c("cornflowerblue","orange")) +
  labs(x = "Sample means", y = "Samples",
       fill = "Size of each sample",
       title = "Distribution of results ('sampling distribution')",
       caption = paste0("Population mean ('true value') = ",round(popmean, digits = 2),
                      "\n Mean over all results ('mean of means') for samples of 1000 = ",
                      round(sampdistmean, digits = 2)," (SD = ",round(sampdist_SD, digits =
                      "\n Mean over all results ('mean of means') for samples of 200 = ",
                      round(sampdist_200mean, digits = 2)," (SD = ",round(sampdist_200_SD,
                      theme(legend.position = "bottom")
```

`stat_bin()` using `bins = 30`. Pick better value `binwidth`.



####... and taking time

How long did all of this take (on my rickety old Intel-based MacBook Pro):

```
# Taking time
end <- Sys.time()
duration <- end - start
duration
```

Time difference of 2.591034 mins